

Ethical Considerations in Deep Learning

Project 4

Name Syahrul Sardonyx Nalli

Registration No. 2411021240061

Subject Deep Learning

Project Project 4 — Ethical Considerations in Deep Learning

INTRODUCTION

Deep Learning is now used in many areas where computer systems have to recognize patterns and make predictions. Examples include image recognition, recommendation systems, healthcare applications, financial services, education, and recruitment. These systems can be useful, but their use also raises ethical questions. A model may give accurate results and still cause problems if the data used to train it is biased, if personal information is not protected, or if people cannot understand how an important decision was made. For this reason, ethical issues should be considered when a Deep Learning system is designed, trained, tested, and used.

TASK 1 — THREE ETHICAL ISSUES IN DEEP LEARNING

1. Bias and Fairness

Bias is one of the main ethical problems in Deep Learning. A model learns from the examples given to it, so problems in the training data can become problems in the final system. If some groups are underrepresented, or if the historical decisions in the data are already unfair, the model can learn those patterns. This can lead to different levels of accuracy or different outcomes for different groups. The issue is especially serious in areas such as hiring, lending, healthcare, and facial recognition because an incorrect or biased prediction can directly affect a person's opportunities. Fairness therefore requires developers to check the data and test the model on different groups rather than looking only at one overall accuracy value.

2. Privacy

Deep Learning systems often need large amounts of data for training. Some of this data can contain personal information, such as photographs, location details, financial records, messages, or health information. Collecting or using this data without proper permission can create privacy problems. Poor security can also expose information to people who should not have access to it. Developers should collect only the information that is needed, protect stored data, control who can access it, and explain clearly how the information will be used. Privacy should be considered from the beginning of a project instead of being treated as an issue only after a system has been built.

3. Transparency and Explainability

Deep Learning models can be difficult to understand because they may contain many layers and a very large number of parameters. This can make it hard for a person to know why a particular prediction was produced. This is often described as the black-box problem. It becomes more important when an AI system is involved in decisions about employment, education, healthcare, loans, or other important parts of a person's life. Users should know when AI is being used and should have enough information to understand the purpose and limitations of the system. Explainability methods, documentation, testing, and human review can make these systems easier to question and monitor.

TASK 2 — REAL-WORLD EXAMPLE: AMAZON'S AI RECRUITING SYSTEM

A well-known example of an ethical concern involved an experimental recruiting system developed by Amazon. The system was intended to help evaluate job applicants by learning from historical résumés and hiring patterns. Reports about the system found that it could disadvantage women because the historical data reflected a technology workforce in which men were more strongly represented. The system therefore risked learning patterns from the past rather than learning a fair standard for selecting candidates.

The important ethical issue in this example is that the model could reproduce an existing imbalance. Even if the system was designed to make recruitment faster, its recommendations could affect whether people received job opportunities. This shows that accuracy or automation alone is not enough to decide whether an AI system is acceptable. Developers also need to ask where the training data came from, what patterns the model may be learning, and whether the results are fair for the people affected by

them. It also shows why automated systems used for important decisions need regular evaluation rather than being trusted simply because they are based on data.

TASK 3 — MEASURES TO REDUCE OR ADDRESS THE ISSUE

The risk of bias can be reduced through several steps. The first step is to examine the training data before using it. Developers should check whether important groups are represented and whether labels or historical decisions contain unfair patterns. Data preparation should be documented so that it is clear where the information came from and what limitations it has.

The model should then be tested across relevant groups instead of relying only on overall accuracy. If the system performs well for one group but poorly for another, that difference needs to be investigated. Fairness checks should also continue after deployment because a model may behave differently when it receives new data. Regular audits can help identify changes or problems early.

Human oversight is another important measure, especially when an AI prediction can have a major effect on a person. A human reviewer should be able to question or reject an automated recommendation when there is evidence that it is wrong or unfair. Organizations should also document the purpose of the model, its data sources, its limitations, and its testing results. These steps do not guarantee that a system will never be biased, but they make problems easier to find and correct.

CONCLUSION

Ethical considerations are an important part of developing and using Deep Learning systems. Bias and fairness, privacy, and transparency can affect whether an AI system is safe and acceptable in real situations. The example of Amazon's recruiting system shows how historical data can cause an automated system to repeat existing patterns of inequality. Careful data checking, fairness testing, privacy protection, clear documentation, regular audits, and human oversight can help reduce these risks. The goal is not only to build models that work, but also to use them in a way that is fair, responsible, and understandable.